

Joining Data Tables

Motivation

An online retailer stores customer data in two places:

customers

| id | name |
|----|---------|
| 4 | Tukey |
| 8 | Wickham |
| 15 | Mason |
| 16 | Jordan |
| 23 | Patil |
| 42 | Cox |

orders

| order | id | date |
|-------|----|--------|
| 1 | 4 | 1-Jan |
| 2 | 8 | 1-Feb |
| 3 | 42 | 15-Apr |
| 4 | 50 | 17-Apr |

Types of joins

1. **Mutating joins** - add new variables to one data frame from matching observations in another
2. **Filtering joins** - filter observations from one data frame based on whether or not they match an observation in the other


Keys

Used to connect two data tables

- **primary key** uniquely identifies an observation in its own table
- **foreign key** uniquely identifies an observation in another table

inner_join

An inner_join matches pairs of observations when their keys are equal



inner_join

```
inner_join(x = orders, y = customers, by = "id")
```

customers

| id | name |
|----|---------|
| 4 | Tukey |
| 8 | Wickham |
| 15 | Mason |
| 16 | Jordan |
| 23 | Patil |
| 42 | Cox |

orders

| order | id | date |
|-------|----|--------|
| 1 | 4 | 1-Jan |
| 2 | 8 | 1-Feb |
| 3 | 42 | 15-Apr |
| 4 | 50 | 17-Apr |




| order | id | date | name |
|-------|----|--------|---------|
| 1 | 4 | 1-Jan | Tukey |
| 2 | 8 | 1-Feb | Wickham |
| 3 | 42 | 15-Apr | Cox |

Outer joins

Keep the rows that appear in a specified table


Left



| key | val_x | val_y |
|-----|-------|-------|
| 1 | x1 | y1 |
| 2 | x2 | y2 |
| 3 | x3 | NA |

A **left join** keeps all observations in the x table


Right



| key | val_x | val_y |
|-----|-------|-------|
| 1 | x1 | y1 |
| 2 | x2 | y2 |
| 4 | NA | y3 |

A **right join** keeps all observations in the y table

Full



| key | val_x | val_y |
|-----|-------|-------|
| 1 | x1 | y1 |
| 2 | x2 | y2 |
| 3 | x3 | NA |
| 4 | NA | y3 |

A **full join** keeps all observations in both the x and y tables


```
left_join(x = orders, y = customers, by = "id")
```

customers

| id | name |
|----|---------|
| 4 | Tukey |
| 8 | Wickham |
| 15 | Mason |
| 16 | Jordan |
| 23 | Patil |
| 42 | Cox |

orders

| order | id | date |
|-------|----|--------|
| 1 | 4 | 1-Jan |
| 2 | 8 | 1-Feb |
| 3 | 42 | 15-Apr |
| 4 | 50 | 17-Apr |



| order | id | date | name |
|-------|----|--------|---------|
| 1 | 4 | 1-Jan | Tukey |
| 2 | 8 | 1-Feb | Wickham |
| 3 | 42 | 15-Apr | Cox |
| 4 | 50 | 17-Apr | NA |

`right_join(x = orders, y = customers, by = "id")`

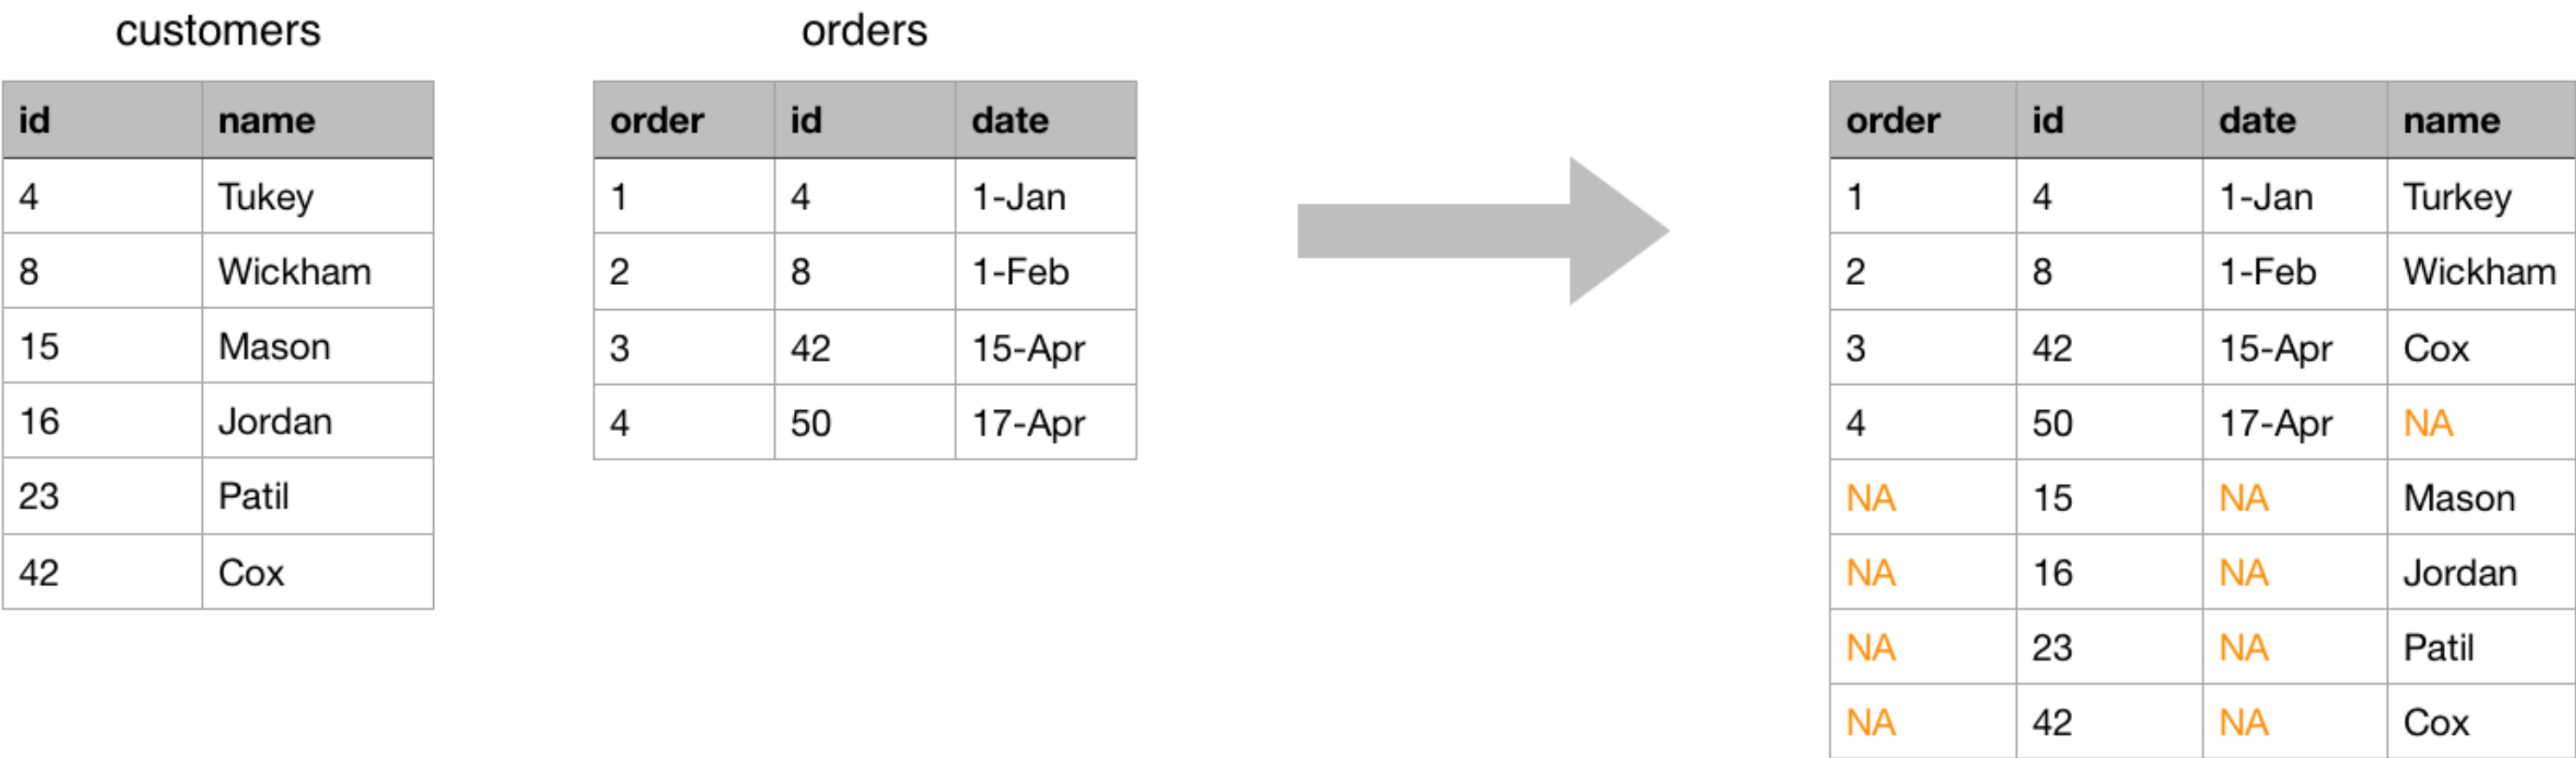
| customers | |
|-----------|---------|
| id | name |
| 4 | Tukey |
| 8 | Wickham |
| 15 | Mason |
| 16 | Jordan |
| 23 | Patil |
| 42 | Cox |

| orders | | |
|--------|----|--------|
| order | id | date |
| 1 | 4 | 1-Jan |
| 2 | 8 | 1-Feb |
| 3 | 42 | 15-Apr |
| 4 | 50 | 17-Apr |



| order | id | date | name |
|-------|----|--------|---------|
| 1 | 4 | 1-Jan | Tukey |
| 2 | 8 | 1-Feb | Wickham |
| NA | 15 | NA | Mason |
| NA | 16 | NA | Jordan |
| NA | 23 | NA | Patil |
| 3 | 42 | 15-Apr | Cox |

```
full_join(x = orders, y = customers, by = "id")
```




| order | id | date | name |
|-------|----|--------|---------|
| 1 | 4 | 1-Jan | Tukey |
| 2 | 8 | 1-Feb | Wickham |
| 3 | 42 | 15-Apr | Cox |
| 4 | 50 | 17-Apr | NA |
| NA | 15 | NA | Mason |
| NA | 16 | NA | Jordan |
| NA | 23 | NA | Patil |
| NA | 42 | NA | Cox |

Filtering joins


Filtering joins still match observations between two data tables, but do not add additional variables, they **only impact the rows returned**.

1. `semi_join(x, y)` keeps all observations in `x` that have a match in `y`
2. `anti_join(x, y)` drops all observations in `x` that have a match in `y`


semi_join(x, y)



Observations will never be duplicated.



anti_join(x, y)



| key | val_x |
|-----|-------|
| 3 | x3 |

```
semi_join(x = orders, y = customers, by = "id")
```

orders

| order | id | date |
|-------|----|--------|
| 1 | 4 | 1-Jan |
| 2 | 8 | 1-Feb |
| 3 | 42 | 15-Apr |
| 4 | 50 | 17-Apr |

customers

| id | name |
|----|---------|
| 4 | Tukey |
| 8 | Wickham |
| 15 | Mason |
| 16 | Jordan |
| 23 | Patil |
| 42 | Cox |



| order | id | date |
|-------|----|--------|
| 1 | 4 | 1-Jan |
| 2 | 8 | 1-Feb |
| 3 | 42 | 15-Apr |

What if we had an extra order?

```
extra_order <- data.frame(order = 5, id = 42, date = "May-01")  
orders2 <- rbind(orders, extra_order)
```

customers

| id | name |
|----|---------|
| 4 | Tukey |
| 8 | Wickham |
| 15 | Mason |
| 16 | Jordan |
| 23 | Patil |
| 42 | Cox |

orders2

| order | id | date |
|-------|----|--------|
| 1 | 4 | 1-Jan |
| 2 | 8 | 1-Feb |
| 3 | 42 | 15-Apr |
| 4 | 50 | 17-Apr |
| 5 | 42 | 1-May |

How do semi_join and inner_join differ?

```
semi_join(x = customers, y = orders2, by = "id")
```

customers

| id | name |
|----|---------|
| 4 | Tukey |
| 8 | Wickham |
| 15 | Mason |
| 16 | Jordan |
| 23 | Patil |
| 42 | Cox |

orders2

| order | id | date |
|-------|----|--------|
| 1 | 4 | 1-Jan |
| 2 | 8 | 1-Feb |
| 3 | 42 | 15-Apr |
| 4 | 50 | 17-Apr |
| 5 | 42 | 1-May |



| id | name |
|----|---------|
| 4 | Tukey |
| 8 | Wickham |
| 42 | Cox |

How do semi_join and inner_join differ?

```
inner_join(x = customers, y = orders2, by = "id")
```

customers

| id | name |
|----|---------|
| 4 | Tukey |
| 8 | Wickham |
| 15 | Mason |
| 16 | Jordan |
| 23 | Patil |
| 42 | Cox |

orders2

| order | id | date |
|-------|----|--------|
| 1 | 4 | 1-Jan |
| 2 | 8 | 1-Feb |
| 3 | 42 | 15-Apr |
| 4 | 50 | 17-Apr |
| 5 | 42 | 1-May |



| id | date | name |
|----|--------|---------|
| 4 | 1-Jan | Tukey |
| 8 | 1-Feb | Wickham |
| 42 | 15-Apr | Cox |
| 42 | 1-May | Cox |

For an anti_join, order matters

```
anti_join(x = orders2, y = customers, by = "id")
```

customers

| id | name |
|----|---------|
| 4 | Tukey |
| 8 | Wickham |
| 15 | Mason |
| 16 | Jordan |
| 23 | Patil |
| 42 | Cox |

orders2

| order | id | date |
|-------|----|--------|
| 1 | 4 | 1-Jan |
| 2 | 8 | 1-Feb |
| 3 | 42 | 15-Apr |
| 4 | 50 | 17-Apr |
| 5 | 42 | 1-May |



| order | id | date |
|-------|----|--------|
| 4 | 50 | 17-Apr |

For an anti_join, order matters

```
anti_join(x = customers, y = orders2, by = "id")
```

customers

| id | name |
|----|---------|
| 4 | Tukey |
| 8 | Wickham |
| 15 | Mason |
| 16 | Jordan |
| 23 | Patil |
| 42 | Cox |

orders2

| order | id | date |
|-------|----|--------|
| 1 | 4 | 1-Jan |
| 2 | 8 | 1-Feb |
| 3 | 42 | 15-Apr |
| 4 | 50 | 17-Apr |
| 5 | 42 | 1-May |



| id | name |
|----|--------|
| 23 | Patil |
| 16 | Jordan |
| 15 | Cox |

Common complications

Joining by multiple variables,

You must specify a vector of variable names:

```
by = c("var1", "var2", "var3").
```

All three columns must match in both tables.

Use all variables that appear in both tables

Leave the by argument blank.

Column names differ between tables

```
by = c("left_var" = "right_var").
```


Your Turn

Does payroll differ between the American League and the National League?

- Load the `tidyverse`
- Install and load the `Lahman` R package
- Look at the `Salaries` and `Teams` data tables
- Devise a way to clearly compare the team payroll between the two leagues over the years